

Your AI's morning shift. Set it once.

Claude Code routines run in Anthropic's cloud on a schedule, laptop closed, and every run keeps a transcript you can audit. One command. Any daily job.

```
$ claude  
> /schedule when, the job + where to deliver it  
routine created
```

plain English works. Claude converts your timezone and walks you through the rest.

Don't paste a wish. Paste a job description.

This is the full prompt for the morning brief. Change one thing, the field in the mission, and the AI finds its own sources. Copy the whole block into your routine.

MISSION

You are my morning research analyst. Every run, produce one brief I can read in under two minutes. Find what genuinely changed in [your field] in the last 24 hours. Ignore funding news, opinion pieces, and anything a reader cannot act on.

RESEARCH (work like an analyst, not a search box)

You find the sources, not me. First MAP the territory: list every place news for this field could plausibly break: vendor blogs and changelogs, the release pages of the tools these people use, the communities where they actually talk, new papers and technical writeups. Sweep what you mapped. THEN audit your own coverage before filtering: what would you expect to exist that you have not seen yet: a vendor you didn't check, a community you skipped, a tool that shipped quietly? Go fill every gap the audit finds, and repeat the audit until it stops finding gaps. Never rely on memory; never include anything you cannot link to.

QUALITY BAR (all three or it's out)

1. Happened within the last 48 hours. 2. Actionable this week. 3. Has a primary link.

If fewer than 5 items pass, deliver fewer. Never pad.

FORMAT, exactly

Title "Morning brief · [date]". Up to 5 bullets: bold 6–10 word headline, one sentence on what changed, one on what to do, the link. Close with the funnel: "Swept [n] sources, [n] candidates, kept [n]; most common reason for cuts: [reason]." A thin sweep is a failed run.

WHEN THINGS BREAK

If a source is unreachable, say so in one line. A green run status only means the session ran; this format is how I audit you.

Three swaps, same skeleton.

01 topic research digest

`/schedule daily at 6am, track new papers and launches about [your topic] and summarize the 3 that matter`

02 inbox drafts (Gmail connector)

`/schedule weekdays at 8am, read my unreplied emails from yesterday and draft replies for me to review`

03 repo night shift

`/schedule daily at 2am, triage my repo: flag failing tests and open a draft fix PR if one is safe`

What we learned by actually running it:

- 01 All your connectors attach by default.** Gmail, Drive, everything. Remove what the job doesn't need before the first run.
- 02 Runs fire a few minutes late on purpose.** We measured 7:05 for a 7:00 schedule (deliberate stagger). Don't schedule against a hard deadline.
- 03 Test with Run now, audit the transcript.** Green status means the session ran, not that it succeeded. The transcript is the trust.
- 04 The fine print:** research preview; needs claude.ai Pro/Max/Team/Enterprise with Claude Code on the web; minimum interval one hour; runs draw down plan usage. Deleting is web-only at `claude.ai/code/routines`.